



Introduction to *cgmlib*

School of Computer Science
Carleton University
Ottawa

Presenter: Albert Chan
(achan@scs.carleton.ca)
(<http://www.scs.carleton.ca/~achan>)

2003-01-31



Agenda

- ☺ Introduction
- ☺ Architecture
- ☺ Current Status
- ☺ Features
- ☺ The *cgmlib* in Details
- ☺ Extension
- ☺ Installation
- ☺ Using *cgmlib*
- ☺ Demo
- ☺ Q & A



What is CGM?

- ☺ CGM stands for Coarse Grained Multi-computers
- ☺ Invented by Dr. Frank Dehne et al. (Carleton University)
- ☺ Number of processors significantly less than input data size.
- ☺ Two parameters:
 - n = size of data;
 - p = number of processors.
- ☺ CGM algorithms consist alternating local computation and global communication rounds.
- ☺ In each communication rounds, an h -relation can be sent and received by each processor, here $h=O(n/p)$
- ☺ Efficient CGM algorithms minimize both the local computation time and the number of global communication rounds.



Why CGM?

- ☺ Why coarse grained?
 - It is difficult to build fine grained machine.
 - We have to use existing machines to simulate fine grained algorithms
 - The simulation usually results a large amount of small messages to be sent and received from time to time.
 - Communication is still the bottleneck of the parallel algorithms, so we prefer a small (or fix) amount of large messages to a large amount of small messages.
- ☺ Why CGM?
 - Although equivalently powerful, other models are usually more difficult to analyze.
 - All messages in a global communication round are grouped into a fix amount of h -relations, thus minimizing the overhead of communication.




Why *cgmlib*?

- ☺ Most CGM machines are based on the message passing mechanism to exchange data.
- ☺ Although the message passing platforms have been settled to only a few different standards, they are still incompatible to each other. This forces the CGM algorithm authors to modify their algorithms when they move the algorithms to a different system (or when the system administrator decides to switch to another platform).
- ☺ The *cgmlib* acts as a “*middle man*” sitting between the CGM algorithms and the underlying message passing platform. Thus allowing the CGM algorithms to be moved to different system with no or minimum changes.
- ☺ The *cgmlib* provides handy methods to allow data easily exchanged amount processors.



The Architecture of *cgmlib*

- ☺ The following diagram shows the relationship of *cgmlib* with other computing elements:





Current Status

- ☺ The *cgmlib* is currently in beta release.
- ☺ It has been installed on the ultra network and the sigma network in the School of Computer Science at Carleton University.
 - ➔ On the ultra network, it is installed at `/home/70user2/cgm/cgmlib-ultra`
 - ➔ On the sigma network, it is installed at `/home/70user2/cgm/cgmlib-sigma`
- ☺ It will also be installed on the thog network in HPCVL at Carleton University and the CGM1 network at Dalhousie University, location to be announced later.
- ☺ You are also encouraged to installed it on your own Linux box or in your UNIX account. See later slide for details on installation.



Features of *cgmlib*

- ☺ Implemented in C++.
- ☺ Completely object-oriented design.
- ☺ Easily extensible.
- ☺ Designed with the mind to promote software reuse.
- ☺ Complete source code available.
- ☺ Many examples included.
- ☺ Comes with on-line user manual.
- ☺ Comes with printable user manual in postscript and PDF formats.
- ☺ Currently only MPI is supported.
- ☺ Provides facilities for high level operations such as sorting, prefix sum calculation and array balancing.
- ☺ Allows the CGM to be partitioned into a group of smaller CGMs.



What can be sent?

- ☺ Everything you want to send across the network **MUST** be sub-classed from interface CommObject, which defines the following methods:
 - copyToArray;
 - copyFromArray;
 - clone;
 - operator=;
 - sendToOstream;
 - getSize;
- ☺ Note that most of the methods are pure virtual so you must provide a **CORRECT** implementation to these methods before you can really use your objects.
- ☺ The only exception is “sendToOstream” which, unless overridden, defaults to print out its memory address.



But I Only Want to Send a Simple Integer...

- ☺ A class called BasicCommObject is defined to allow the user to send simple data type (int, long, float, double, char, etc).
- ☺ To use BasicCommObject, do the following:
 - CommObject *intData = new BasicCommObject <int> (3);
 - CommObject *floatData = new BasicCommObject <float> (5.5);
 - CommObject *charData = new BasicCommObject <char> ('a');
- ☺ You cannot use BasicCommObject on pointer types as it is meaningless to send pointers to other processors:
 - pointer is only an address to the data; and
 - each processor has its own address space.
- ☹ Never try to do the following (it won't work):
 - ✗ CommObject *pointerData = new BasicCommObject (char *) ("Hello");



Can I Send a Non CommObject?

- ☺ Yes, a class called SimpleCommObject is defined to allow the user to wrap a non-CommObject into a CommObject.
- ☺ To use SimpleCommObject, do the following:
 - ➔ `CommObject *myData = new SimpleCommObject <MyDataType> (x);`
 - ➔ Here x is an object of MyDataType.
 - ➔ The object to be wrap around into a SimpleCommObject **CANNOT** contain any pointer.



CommObjectList

- ☺ A container to encapsulate the CommObject's to be sent across the network.
- ☺ The data type that most of the methods in *cgmlib* expect.
- ☺ Easier to use than CommObject ** as the container will manage its own memory.
- ☺ Automatically expandable.
- ☺ Array-like behavior.
 - ➔ `CommObjectList data;`
 - `CommObject *firstElement = data [0];`
- ☺ Built-in boundary check.



Comm

- ☺ The heart of *cgmlib*.
- ☺ Interface only, that means all methods are pure virtual.
- ☺ Some handy constants defined:
 - ANY_NODE = -1;
 - PROC_ZERO = 0;
 - MPICOMM="MPICOMM";
- ☺ A static method is provided to encapsulate the underlying message passing platform:
 - getComm;



Comm

- ☺ Other (non-static) methods that Comm provides:

| | |
|--------------------------|------------------|
| → getNumberOfProcessors; | → allToAllBCast; |
| → getMyId; | → hRelation; |
| → synchronize; | → arrayBalancing |
| → send; | → partitionCGM |
| → receive; | → unPartitionCGM |
| → oneToAllBCast; | → dispose |
| → allToOneGather; | |
- ☺ A concrete implementation is provided in the following class:
 - MPIComm



Request System

- ☺ A request system has been implemented to help individual elements requesting information across the network.
- ☺ Use consolidation technique to avoid overloading any particular processor.
- ☺ Requests sent to the same elements are consolidated into a set of at most p requests.
- ☺ Each processor is handling at most $n/p + p$ requests.
- ☺ Responses are automatically expanded and routed to the original requestor.
- ☺ Each requestor will get its own copy of the response.



Other Operations

- ☺ The following classes provide other higher level operations:
- ☺ ParallelSorter
 - sort
- ☺ ParallelPrefixSummer
 - calculatePrefixSum



Other Utilities

- ☺ Some other utilities are also provided :
 - Timer: for getting timing information
 - ✧ UnixTimer: use Unix wall clock
 - ✧ CPUTimer: use CPU ticks
 - ✧ CGMTimers: a set of timers to time:
 - total elapse time;
 - total communication time; and
 - total computational time.
 - Random: for generating random numbers
 - Sorter: for sorting CommObjects
 - ✧ HeapSorter: use heap sort - $O(n \log n)$;
 - ✧ IntegerSorter: use integer sort - $O(n)$.
 - GeneralUtilities
 - ✧ Provide miscellaneous utilities that may be useful for the users.



Extending *cgmlib*

- ☺ *cgmlib* can be easily extended to provide support to other message passing platforms, such as PVM.
- ☺ In addition to provide a concrete extension of the Comm class (e.g. PVMComm), you also need to modify the following *IN* the Comm Class:
 - add an include statement to include your new class
 - ✧ e.g. `#include "PVMComm.h"`
 - add an interface string
 - ✧ e.g. `static char *PVMCOMM = "PVMCOMM";`
 - change the default if desired
 - ✧ e.g. `static char *defaultComm = PVMCOMM;`
 - add another "else if" statement in the getComm method
 - ✧ see next slide for example.



Extending cgmlib

```
Comm *Comm::getComm (int *argc, char ***argv, CGMTimers *timers, char *selector)
{
    if (selector == NULL)
    {
        selector = defaultComm;
    }
    if (strcmp (selector, MPICOMM) == 0)
    {
        return MPIComm::getComm (argc, argv, timers);
    }
    else if (strcmp (selector, PVMCOMM) == 0)
    {
        return PVMComm::getComm (argc, argv, timers);
    }
    return getComm (argc, argv, timers, defaultComm);
}
```



Installation

- ☺ The library is distributed under Gnu's LGPL:
→ <http://www.gnu.org>
- ☺ Due to the diversity of the parallel platforms, only source code is distributed.
- ☺ A WWW page has been set up to be the home of *cgmlib*:
→ <http://www.scs.carleton.ca/~cgm>
- ☺ Down load the file *cgmlib.src.zip* or *cgmlib.src.tar.gz*
- ☺ Unzip the file
- ☺ Modify the *Makefile* and *Makefile.common* if necessary.
- ☺ Execute "make" (on the top level directory of *cgmlib*)
- ☺ If everything is OK, you can install it by executing "make install"



Using *cgmlib*

- ☺ Assume *cgmlib* is installed at `/usr/local/cgmlib`
- ☺ To compile a file:
 - `g++ -I/usr/local/cgmlib/include -c <file.cpp>`
- ☺ To link an object module:
 - `g++ -o <executable> <file.o> -L/usr/local/cgmlib/lib -lsgm <otherflag>`
 - here `<otherflag>` is the flag that is specific to the message passing platform being used
 - for example, for MPICH, it may be
 - ✧ `-L/usr/local/mpich/lib -lmpich++ -lmpich`
- ☺ To run the linked program, just follow the normal procedure in your environment, e.g.:
 - `mpirun -np <#> <program> <args>`



Application: *cgmgraph*

- ☺ Several examples (mostly graph applications) have been provided, and grouped as *cgmgraph*:
 - List Ranking
 - Euler Tours
 - Connected Components
 - Spanning Forest
 - Bipartite Graph Detection



Future Work

- ☺ Provide error handling.
- ☺ Provide support for other message passing platforms such as PVM.
- ☺ Provide Java binding (e.g. through mpiJava or Java RMI)
- ☺ Provide automatic installation using GNU automake and autoconf.
- ☺ Provide default implementation to “middle-level” operation such as array-balancing, broadcast and gather.
- ☺ Integrate *cgmgraph* into LEDA.
- ☺ Provide implementation of more graph algorithms
- ☺ Performance tuning.



Demo

